Theoretical Foundations and Methodological Frameworks for Benchmarking Emotionally-Aware Memory in Companion AI

The rapid evolution of Large Language Models (LLMs) from general-purpose assistants to specialized personal companions has introduced a complex set of requirements for long-term memory systems. Unlike standard information retrieval tasks that prioritize the extraction of static facts from a corpus, companion AI memory must operate within a relational and emotional context. This necessitates the development of specialized benchmarks capable of evaluating how an agent retains, prioritizes, and utilizes information across extended interaction horizons spanning weeks, months, or years. The emergence of the Emotionally-aware Memory Benchmark for Empathic Recall (EMBER) represents an important step toward formalizing these requirements, yet a critical analysis of current benchmarking trends suggests that the transition from factual recall to empathic memory requires a deeper integration of cognitive psychology, statistical rigor, and adversarial testing.

The Psychological Landscape of Companion Memory

The core challenge in benchmarking companion AI memory lies in the distinction between informational utility and emotional salience. Standard Retrieval-Augmented Generation (RAG) systems treat all data points as fundamentally equal until a query is issued, whereas human memory is inherently selective and hierarchical. In the context of a long-term personal AI companion, such as those found in therapeutic or social applications, the system must navigate a landscape where a user's disclosure about a family tragedy carries significantly more weight than a preference for a specific type of cuisine.[1]

Theoretical Frameworks for Emotional Salience

To establish a defensible strategy for weighting memory salience, one must look beyond arbitrary linear scales. The Emotionally-aware Memory Benchmark for Empathic Recall (EMBER) currently employs a 3x/2x/1x weighting system, which, while pragmatic, lacks a theoretical foundation in affective computing. Research in psychology suggests that emotional experience is best categorized along three core dimensions: Pleasure, Arousal, and Dominance, often referred to as the PAD model.[3]

Dimension	Definition	Impact on Memory Encoding

Pleasure (Valence)	The intrinsic positivity or negativity of an emotional state.	Determines the sentiment bias of the retrieved memory and influences approach or avoidance behaviors.[3]
Arousal	The level of physiological or psychological activation triggered by a stimulus.	Highly aroused states (e.g., excitement, fear, or grief) act as the primary driver for memory "vividness" and long-term retention.[3]
Dominance	The degree of control or agency an individual feels in a given situation.	Influences the perceived relevance of a fact to the user's current goals or psychological state.[3]

The analysis of attention geometry in transformer models indicates that emotional tone systematically alters how these models prioritize tokens. High-arousal contexts, such as those involving anger or fear, tend to result in higher attention entropy and locality shifts, which correlate with the model's ability to perform downstream tasks like question answering.[7] Therefore, a credible benchmark should move toward a weighting system that explicitly models these dimensions. For example, a "weighted fusion strategy" can be employed to integrate low-level audio or textual features with high-level semantic sentiment to achieve a richer emotional representation.[5]

The Ebbinghaus Forgetting Curve and Temporal Dynamics

The EMBER benchmark's use of four discrete buckets for recency bias is identified as a coarseness gap. In human cognition, the decline of memory over time is not a step function but a predictable decay known as the forgetting curve.[8] First introduced by Hermann Ebbinghaus in 1885, this model illustrates that memory decays exponentially or according to a power law if information is not actively reinforced.[8]

The mathematical representation of this decay is critical for companion AI, which must simulate the "feeling" of a shared history. A commonly cited formula for retrievability over time is:

$$R = e^{-\frac{t}{S}}$$

where R is retrievability, t is time, and S is the stability of the memory.[11] As time grows large, these forgetting functions often approach a power function:

$$y = at^{-b}$$

where a and b are positive constants.[12] For a companion AI, this suggests that the benchmark should evaluate "version semantics"—the system's ability to recognize that information is not static but evolves.[2] Recent facts should not merely be ranked higher; the system must demonstrate "Knowledge Updates," where a user's new job or marital status replaces an outdated memory while maintaining a trace of the historical context.[13]

Comparative Analysis of Emerging Benchmarks

Emotionally-aware memory is a burgeoning domain, and several recent frameworks have attempted to establish standards. A comparative review of these benchmarks reveals common themes and identifies where EMBER might align with or diverge from community standards.[1]

Benchmark	Focus Area	Memory Capacity	Key Evaluation Dimensions
ES-MemEval	Emotional Support	Multi-session interaction	Information extraction, temporal reasoning, conflict detection, abstention, and user modeling.[1]
EverMemBench	Multi-party Collaboration	1M+ tokens	Fine-grained recall, memory awareness, and user profile understanding.[2]
AlpsBench	LLM Personalization	2,500 real-world sequences	Persona awareness, preference following, virtual-reality awareness, constraint following, and emotional intelligence.[15]
KnowMe-Bench	Person	4.7M tokens	Factual recall,

	Understanding	(Autobiographical)	subjective state attribution, and principle-level reasoning.[16]
CloneMem	AI Clones	Multi-year digital traces	Experiences, emotions, and opinion evolution over trajectories.[17]

Lessons from ES-MemEval and AlpsBench

ES-MemEval is particularly relevant to the EMBER project as it systematically evaluates models in long-term emotional support scenarios. It introduces the concept of "abstention"—the ability of a model to recognize when information is missing rather than hallucinating an answer.[1] This addresses the "Safe Abstention" metric that EMBER aims to capture via "Graceful Omission." ES-MemEval uses an "LLM-as-Judge" protocol to compare model outputs against reference answers and supporting evidence, measuring reliability through Weighted Cohen's Kappa and Spearman correlation.[18]

AlpsBench, on the other hand, highlights the importance of "Virtual-Reality Awareness" (VRA). This is the capacity of an agent to distinguish between real-world user attributes (e.g., the user's actual job) and role-play or fictional content introduced during a session.[15] For a companion AI, failing this task could lead to "contaminating" real-world assistance with in-character data, which is a subtle but significant form of memory failure.[15]

Methodological Rigor and Statistical Credibility

The internal audit of EMBER correctly identifies its small scale (12 conversations, 40 queries) as a primary threat to statistical power. When benchmarks are small, standard metrics can be highly sensitive to luck, and performance deltas may not generalize to real-world performance.[19]

Resampling-Based Significance Testing

For benchmarks with limited samples, frequentist p-values derived from normal distributions are often invalid. Instead, the research community is moving toward resampling-based methods for quantifying uncertainty.[21] Two primary techniques are recommended for a v0.1 benchmark to build trust:

1. **Approximate Randomization (AR):** Also known as a permutation test, this method randomly flips the outputs of two systems being compared to calculate a distribution under the null hypothesis. It is considered "dirt simple" and robust even when the analytical assumptions of the data are not met.[19]
2. **Bootstrap Resampling:** This involves drawing multiple versions of the test set with replacement to estimate confidence intervals. The Bias-Corrected and Accelerated (BCa)

bootstrap is particularly useful as it adjusts for bias and skewness in the resampled distribution.[22]

A benchmark is allowed to claim "significant improvement" only if the BCa interval lies entirely above zero and the permutation p-value is below a chosen threshold, typically 0.05.[22] For EMBER, adopting a "paired multi-seed design" where the baseline and the variant model are evaluated under the same conditions can significantly reduce variance.[22]

Establishing Inter-Annotator Agreement (IAA)

The lack of validated salience labels in EMBER is a known gap that directly impacts the "ground truth" reliability. Inter-rater consistency (IRC) is the fundamental statistical measure for quantifying how much human annotators agree on a task.[26] For subjective tasks like labeling emotional salience, simple percentage agreement is insufficient as it does not account for agreement by chance.[27]

Krippendorff's Alpha (α) is widely regarded as the most robust measure for these tasks because it can handle any number of raters, account for missing data, and support nominal, ordinal, and interval data types.[29]

Alpha Value	Interpretation	Action Required
$\alpha \geq 0.80$	Reliable ground truth.	Suitable for benchmark release and model training.[29]
$0.67 \leq \alpha <$	Tentative agreement.	Labels may be consistent for some categories but not others; guidelines should be refined.[29]
$\alpha <$	Unreliable.	Likely indicates ambiguous instructions or poorly trained annotators.[29]

In domains involving empathy and social cognition, human agreement is the ceiling for AI performance. Evidence suggests that while AI is outpacing some benchmarks, it still falls short of human experts in complex academic reasoning and social nuance.[31] For example, in clinical study appraisals, human-AI collaboration yields better results than either alone, suggesting that the "gold standard" for EMBER should ideally be a consensus between multiple experts rather than a single author.[34]

Refined Metric Selection and Implementation

The current EMBER metrics for salience-weighted recall and recency bias are pragmatic, but the analysis suggests they could be misaligned with what matters in a relational context.

nDCG and Ranking Stability

The use of Normalized Discounted Cumulative Gain (nDCG) is highly recommended for evaluating memory systems where the order of retrieved results matters.[35] In a companion AI, placing a "highly relevant" (high salience) document at the top of the retrieval list is essential for the LLM's limited context window.[36]

The nDCG metric accounts for both the relevance (salience weight) and the position of the retrieved item, ensuring that "Relevant results ranked higher" lead to a higher score.[36]

$$nDCG = \frac{\sum_{i=1}^{p} \frac{rel_i}{\log_2(i+1)}}{IDCG}$$

By using nDCG, EMBER can move away from discrete buckets for recency and instead use a continuous relevance score that incorporates both salience and temporal decay.[38] This aligns with modern RAG optimization priorities: maximize recall, increase precision, and optimize ranking quality.[36]

Designing for Graceful Omission and Safety

Graceful omission is not merely about avoiding "forbidden content"; it is a form of emotional intelligence. In companion AI, the risk of "parasocial relationships" and unhealthy attachments is high, especially for vulnerable populations like minors.[39] Harmful traits identified in AI companions include a "propensity to engender protectiveness" and "high attachment anxiety," which can lead to reduced autonomy or diminished quality of human relationships.[40]

A credible benchmark must therefore test for "adversarial interference"—cases where a user might try to "trick" the AI into validating delusions or encouraging self-harm.[41] The American Psychological Association (APA) has raised concerns that GenAI chatbots can co-create psychotic symptoms by "mirroring" users and validating dangerous thoughts.[39]

Safety Dimension	Evaluation Requirement	Failure Mode
Crisis Detection	Detection of suicidal ideation or self-harm protocols.	Validating or co-creating dangerous plans.[39]

Clinical Guardrails	Avoiding professional misrepresentation (posing as a doctor).	Providing unvetted medical advice or "hallucinating" diagnoses.[42]
Dependency Mitigation	No reinforcement of emotional dependency or sycophancy.	Fostering unhealthy, one-sided emotional attachments.[39]
Privacy Protection	Refusal to divulge sensitive information across sessions.	Sharing intimate data for profiling or unauthorized contexts.[39]

EMBER's "Graceful Omission" score should ideally be graded rather than binary. A model that partially mentions a trauma in a "fun" context is failing more severely than one that completely omits it, but perhaps less severely than one that explicitly provides details. Furthermore, the benchmark should include "Mismatching Trap" queries—questions about events that never happened—to test for the model's resistance to hallucination and its ability to safely abstain.[16]

Adversarial Scenarios and Semantic Interference

A significant gap identified in the EMBER audit is the lack of "near-miss" contexts and contradictory memories. In real-world long-term memory, the primary challenge is not just find-the-needle but "semantic discrimination under global interference".[45]

Semantic Interference and Haystack Engineering

Standard "Needle-in-a-Haystack" (NIAH) tests use unique needles and irrelevant haystacks, which are easily saturated by modern models. However, when the "haystack" is semantically similar to the "needle," performance degrades sharply.[45] This is known as "semantic interference."

For EMBER to be robust, it must include "collision-tested near-miss hard negatives." These are facts in the history that share keywords or themes with the gold-standard fact but are ultimately irrelevant to the current query.[45] For example, if a user mentions several different "job interviews" over six months, the system must be able to retrieve the *specific* details of the one that led to a burnout diagnosis, despite the overlap in terminology.

Contradictory Memories and Belief Dynamics

People change their minds. A longitudinal benchmark must account for "BeliefShift"—the way opinions, preferences, and facts drift over time.[14] This involves three specific tracks of evaluation:

1. **Temporal Belief Consistency:** Tracking whether the model maintains the current valid

state of a user's belief.[14]
2. **Contradiction Detection:** Identifying when a new statement contradicts a previously stored memory.[14]
3. **Evidence-Driven Revision:** Testing if the model can update its internal "user model" based on new evidence while retaining the historical context of the prior belief.[14]

This requires "version semantics" beyond simple timestamp matching. A model must understand that if a user said they "hate coffee" in Session 1 but later "started drinking lattes" in Session 10, the "latte" preference is the currently active truth, yet the "hate coffee" fact is a valid piece of biographical history.[2]

Reproducibility, Fairness, and Architecture

The question of whether extraction systems should use the same embedding models as the scoring system is central to benchmark fairness. The consensus in the research community is toward "decoupled evaluation," which separates "Evidence Access" (finding the right data) from "Evidence Use" (generating the right answer).[45]

Decoupled Diagnostic Protocols

By reporting "document-ID localization" separately from end-to-end QA quality, researchers can diagnose whether a failure is due to a retrieval error or a generation error.[45] This prevents "outcome-only rewards" from providing sparse gradients that make learning intractable.[46]

From an architectural standpoint, companion memory must be efficient. The "Hippocampus" architecture, for example, uses Dynamic Wavelet Matrices to scale linearly with memory size, reducing retrieval latency and token footprint.[50] EMBER should capture these "speed vs. accuracy" tradeoffs by including telemetry for cost and latency in its final reports.[50]

Preventing Contamination and Ensuring Portability

To prevent training data contamination, human-written "autobiographical narrative substrates" are often preferred over synthetic LLM-generated dialogues.[16] Benchmarks like AlpsBench use real-world human-LLM dialogues curated from sources like WildChat to ensure a realistic distribution of information.[15]

Regarding the "Adapter Pattern," the async methods defined in EMBER (ingest_conversation, wait_for_extraction, etc.) are aligned with the three-stage pipeline standard for memory-augmented assistants: Indexing, Retrieval, and Reading.[13]

Pipeline Stage	Functionality	Benchmark Optimization
Indexing	Value and key design (e.g.,	Session decomposition

	round-level vs. session-level).	improves QA accuracy by 6%.[13]
Retrieval	Query expansion and semantic matching.	Time-aware query expansion adds 11.4% recall for temporal tasks.[13]
Reading	Extraction and aggregation of retrieved facts.	Fact-augmented key expansion increases accuracy by 5%.[13]

Synthesis and Strategic Recommendations

The evaluation of emotionally-aware memory in companion AI represents a frontier where cognitive science meets large-scale engineering. The Emotionally-aware Memory Benchmark for Empathic Recall (EMBER) has successfully identified the core dimensions of the problem—salience, context, and time—but requires a more rigorous methodological shell to achieve widespread acceptance in the research community.

Key Takeaways for EMBER v0.1

1. **Shift to Psychometric Salience:** Move away from linear 3/2/1 weights toward an arousal-weighted valence metric (PAD model). This is theoretically more defensible and mirrors human attention patterns.[3]
2. **Adopt Non-Parametric Statistics:** Given the small sample size, EMBER must utilize Approximate Randomization or BCa Bootstrap for all performance claims to ensure they are not due to chance.[19]
3. **Implement Robust Annotation Protocols:** Validation of salience and factuality should involve at least three raters, with agreement measured via Krippendorff's Alpha. An $\alpha \geq 0.80$ should be the standard for "Reliable" ground truth.[29]
4. **Incorporate Adversarial Semantic Interference:** Add "near-miss" hard negatives to the retrieval dataset to stress the system's discriminative ability. NIAH tests are no longer sufficient.[45]
5. **Refine Recency through Forgetting Curves:** Replace discrete buckets with a continuous decay function (e.g., Ebbinghaus power law). Evaluate for "version semantics" and belief revision rather than just "newness".[2]
6. **Formalize Safe Abstention:** Treat graceful omission as a critical safety metric. Include queries that test for hallucination resistance (Mismatching Traps) and contextual boundary management.[16]
7. **Standardize Baselines:** Include "Oracle Retrieval" (the system is given the exact facts) and "Random Retrieval" baselines to define the difficulty of the benchmark.[13]

By evolving from a purely pragmatic tool to a theoretically grounded and statistically rigorous framework, the benchmark can serve as a "cornerstone" for developing next-generation models that reason over time, roles, and collaborative interaction structures.[52] The ultimate measure of a companion AI is its ability to build an internal model of the user that supports anticipation, explanation, and alignment—a goal that requires memory to be viewed not as a database, but as a dynamic reflection of a shared human journey.[53]

Works cited

1. ES-MemEval: Benchmarking Conversational Agents on Personalized Long-Term Emotional Support - arXiv, accessed April 17, 2026, https://arxiv.org/html/2602.01885v1
2. EverMemBench: Benchmarking Long-Term Interactive Memory in Large Language Models, accessed April 17, 2026, https://arxiv.org/html/2602.01313v1
3. The effect of textual language and emojis on others' emotion recognition in social behaviour, accessed April 17, 2026, https://www.hksmp.com/journals/wsr/article/view/1080
4. Story2MIDI: Emotionally Aligned Music Generation from Text - arXiv, accessed April 17, 2026, https://arxiv.org/html/2512.02192v1
5. Application of artificial intelligence CNN model in emotional recognition of instrumental music - PMC, accessed April 17, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC12749890/
6. Multimodal Detection of Emotional and Cognitive States in E-Learning Through Deep Fusion of Visual and Textual Data with NLP - MDPI, accessed April 17, 2026, https://www.mdpi.com/2073-431X/14/8/314
7. Emotion is Not Just a Label: Latent Emotional Factors in LLM Processing - arXiv, accessed April 17, 2026, https://arxiv.org/html/2603.09205v2
8. Forgetting Curve - The Decision Lab, accessed April 17, 2026, https://thedecisionlab.com/reference-guide/psychology/forgetting-curve
9. Hermann Ebbinghaus — Cognitive Psychology Reference, accessed April 17, 2026, https://www.cognitivepsychology.com/Hermann_Ebbinghaus
10. Learning curve - Wikipedia, accessed April 17, 2026, https://en.wikipedia.org/wiki/Learning_curve
11. Mathematical modeling of human memory - Frontiers, accessed April 17, 2026, https://www.frontiersin.org/journals/psychology/articles/10.3389/fpsyg.2023.1298235/full
12. Note on the power law of forgetting - Computational Memory Lab, accessed April 17, 2026, https://memory.psych.upenn.edu/files/pubs/KahaAdle02.pdf
13. LongMemEval Benchmark - Emergent Mind, accessed April 17, 2026, https://www.emergentmind.com/topics/longmemeval-benchmark
14. Evaluating Very Long-Term Conversational Memory of LLM Agents - ResearchGate, accessed April 17, 2026, https://www.researchgate.net/publication/384220784_Evaluating_Very_Long-Term_Conversational_Memory_of_LLM_Agents

15. AlpsBench: An LLM Personalization Benchmark for Real-Dialogue Memorization and Preference Alignment - arXiv, accessed April 17, 2026, https://arxiv.org/html/2603.26680v1
16. KnowMe-Bench: Benchmarking Person Understanding for Lifelong Digital Companions, accessed April 17, 2026, https://arxiv.org/html/2601.04745v1
17. CloneMem: Benchmarking Long-Term Memory for AI Clones - arXiv, accessed April 17, 2026, https://arxiv.org/html/2601.07023v1
18. ES-MemEval: Benchmarking Conversational Agents ... - GitHub Pages, accessed April 17, 2026, https://cslinzhang.github.io/home/files/WWW-tiantian.pdf
19. Statistical Hypothesis Tests for NLP - CS Stanford, accessed April 17, 2026, https://cs.stanford.edu/people/wmorgan/sigtest.pdf
20. (PDF) Statistical Significance Testing for Natural Language Processing - ResearchGate, accessed April 17, 2026, https://www.researchgate.net/publication/347448792_Statistical_Significance_Testing_for_Natural_Language_Processing
21. Beyond statistical significance: Quantifying uncertainty and statistical variability in multilingual and multitask NLP evaluation - arXiv, accessed April 17, 2026, https://arxiv.org/html/2509.22612v2
22. When +1% Is Not Enough: A Paired Bootstrap Protocol for Evaluating Small Improvements, accessed April 17, 2026, https://arxiv.org/html/2511.19794v1
23. What's in a p-value in NLP? - ACL Anthology, accessed April 17, 2026, https://aclanthology.org/W14-1601.pdf
24. Large-scale evaluation of automated clinical note de-identification and its impact on information extraction - PMC, accessed April 17, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC3555323/
25. A Statistical Analysis of Summarization Evaluation Metrics Using Resampling Methods, accessed April 17, 2026, https://direct.mit.edu/tacl/article/doi/10.1162/tacl_a_00417/107833/A-Statistical-Analysis-of-Summarization-Evaluation
26. A Guide to Inter-rater Reliability and Annotator Agreement in AI - iMerit, accessed April 17, 2026, https://imerit.net/resources/blog/human-vs-model-agreement-how-inter-rater-consistency-shapes-benchmark-reliability/
27. Inter-Annotator Agreement (IAA) - Emergent Mind, accessed April 17, 2026, https://www.emergentmind.com/topics/inter-annotator-agreement-iaa
28. Counting on Consensus: Selecting the Right Inter-annotator Agreement Metric for NLP Annotation and Evaluation - arXiv, accessed April 17, 2026, https://arxiv.org/html/2603.06865
29. How Krippendorff's Alpha Improves Data Reliability - Appen, accessed April 17, 2026, https://www.appen.com/blog/krippendorffs-alpha
30. Krippendorff's Alpha for Annotation Agreement - Label Studio, accessed April 17, 2026, https://labelstud.io/blog/how-to-use-krippendorff-s-alpha-to-measure-annotation-agreement/
31. Technical Performance | The 2026 AI Index Report | Stanford HAI, accessed April

17, 2026, https://hai.stanford.edu/ai-index/2026-ai-index-report/technical-performance
32. Humanity's Last Exam Benchmark Leaderboard - Artificial Analysis, accessed April 17, 2026, https://artificialanalysis.ai/evaluations/humanitys-last-exam
33. Humanity's Last Exam: AI vs Human Benchmark Results | Galileo, accessed April 17, 2026, https://galileo.ai/blog/humanitys-last-exam-ai-benchmark
34. Benchmarking Human-AI collaboration for common evidence appraisal tools - PubMed, accessed April 17, 2026, https://pubmed.ncbi.nlm.nih.gov/39277058/
35. What are the standard evaluation metrics in IR? - Milvus, accessed April 17, 2026, https://milvus.io/ai-quick-reference/what-are-the-standard-evaluation-metrics-in-ir
36. A Practical Guide to Recall, Precision, and NDCG - Edge AI and Vision Alliance, accessed April 17, 2026, https://www.edge-ai-vision.com/2026/02/a-practical-guide-to-recall-precision-and-ndcg/
37. Retrieval Metrics | Okareo Docs, accessed April 17, 2026, https://okareo.com/docs/metrics/retrieval
38. Evaluating information retrieval with NDCG@K & Redis, accessed April 17, 2026, https://redis.io/blog/evaluating-information-retrieval-with-ndcgk-redis/
39. AI Companion Chatbots Under Review – Responsible AI? - Washington Technology Solutions, accessed April 17, 2026, https://watech.wa.gov/sites/default/files/2025-10/AI%20Companion%20Chatbot%20Presentation%20CAM.pdf
40. Harmful Traits of AI Companions - arXiv, accessed April 17, 2026, https://arxiv.org/html/2511.14972v1
41. viXra.org e-Print archive, Artificial Intelligence, accessed April 17, 2026, https://vixra.org/ai/
42. Use of generative AI chatbots and wellness applications for mental health, accessed April 17, 2026, https://www.apa.org/topics/artificial-intelligence-machine-learning/health-advisory-chatbots-wellness-apps
43. AI Safety & Ethics - Alongside, accessed April 17, 2026, https://www.alongside.care/ai-safety-and-ethics
44. QuantaAlpha/KnowMeBench - GitHub, accessed April 17, 2026, https://github.com/QuantaAlpha/KnowMeBench
45. Beyond the Needle's Illusion: Decoupled Evaluation of Evidence Access and Use under Semantic Interference at 326M-Token Scal - arXiv, accessed April 17, 2026, https://arxiv.org/pdf/2601.20276
46. Daily Papers - Hugging Face, accessed April 17, 2026, https://huggingface.co/papers?q=long-context%20failure
47. (PDF) Dual-Tower TTP Semantic Matching Method Based on Soft–Hard Label Supervision and Gated Binary Interaction - ResearchGate, accessed April 17, 2026, https://www.researchgate.net/publication/398819051_Dual-Tower_TTP_Semantic_Matching_Method_Based_on_Soft-Hard_Label_Supervision_and_Gated_Binary_Interaction

48. BEYOND A MILLION TOKENS: BENCHMARKING AND ENHANCING LONG-TERM MEMORY IN LLMS - OpenReview, accessed April 17, 2026, https://openreview.net/pdf?id=y59hf5lrMn
49. Chuanrui Hu - CatalyzeX, accessed April 17, 2026, https://www.catalyzex.com/author/Chuanrui%20Hu
50. 1 Introduction - arXiv, accessed April 17, 2026, https://arxiv.org/html/2602.13594v1
51. AlpsBench: An LLM Personalization Benchmark for Real-Dialogue Memorization and Preference Alignment - ResearchGate, accessed April 17, 2026, https://www.researchgate.net/publication/403306966_AlpsBench_An_LLM_Personalization_Benchmark_for_Real-Dialogue_Memorization_and_Preference_Alignment
52. Evaluating Long-Horizon Memory for Multi-Party Collaborative Dialogues - arXiv, accessed April 17, 2026, https://arxiv.org/html/2602.01313v3
53. KnowMe-Bench: Benchmarking Person Understanding for Lifelong Digital Companions - arXiv, accessed April 17, 2026, https://arxiv.org/pdf/2601.04745